

# Homework2

Kai Wang

2018/1/30

## 1 Classifiers for Basketball Courts

a


```
#observed shots
shots = data.frame(x1=c(.75,.85,.85,.15,.05,.05,.85),
                  x2=c(.10,.80,.95,.10,.25,.50,.25),
                  y=c(-1,-1,1,-1,1,1,-1))
shots.data = shots[,c(1,2)] %>% as.matrix()
shots.label = shots[,3] %>% as.matrix()

#perceptron algorithm from HW1
perceptron = function(x, y, epoch) {
  # initialize weight vector
  weight = rep(0, dim(x)[2])
  result = matrix(0, nrow = epoch, ncol = dim(x)[2])
  for (i in 1:epoch) {
    for (j in 1:length(y)) {
      z = sum(weight*x[j, ])
      if(z <= 0) {
        ypred = -1
      } else {
        ypred = 1
      }

      # Update weight
      if (y[j] != ypred) {
        weight = weight + y[j] * x[j,]
      }
    }


    #save weight vectors for each step
    result[i,] = weight
  }
  return(result)
}

weight = perceptron(shots.data, shots.label,100)
pred.train = t(shots.data %*% t(weight))
pred.train[pred.train>0] = 1
pred.train[pred.train<=0] = -1
accuracy = apply(pred.train, 1, function(x) {return(sum(x==shots.label)/length(shots.label))})
accuracy = c(0,accuracy)
#plot accuracy vs. epoch
plot(x = 0:100, y=accuracy, ylim = range(0,1), type = "l", col = "blue", xlab = "epoch", ylab = "accuracy")
```



From the epoch vs. accuracy plot, we can see the perceptron converges at 37th iteration. Since the accuracy is 100%, there's no empirical error for this classifier. We can come up with other linear classifiers which will give the same error(0). For example, an boolean function  $y = f(x_1, x_2) = \mathbb{I}_{-x_1 + x_2 > 0}$ . More generally, as long as the slope of the linear classifier passing through origin is between  $(\frac{16}{17}, \frac{19}{17})$ .

```
#plot observed data and another seperation line
ggplot(shots) + geom_point(aes(x = x1, y = x2, col = as.factor(y))) +
  theme_bw() + geom_abline(slope = -weight[100,1]/weight[100,2], color = 'red') +
  geom_abline(slope = 1, intercept = 0, color = 'blue') +
  stat_function(fun=function(x) sqrt(x), xlim = c(0,1))
```



The above plot shows data and the decision boundary of the perceptron (as red line). In addition, the boolean function  $y = f(x_1, x_2) = \mathbb{I}_{-x_1+x_2+0.08 > 0}$  is also plotted as a black line, which clearly separates our data with no empirical error.

b

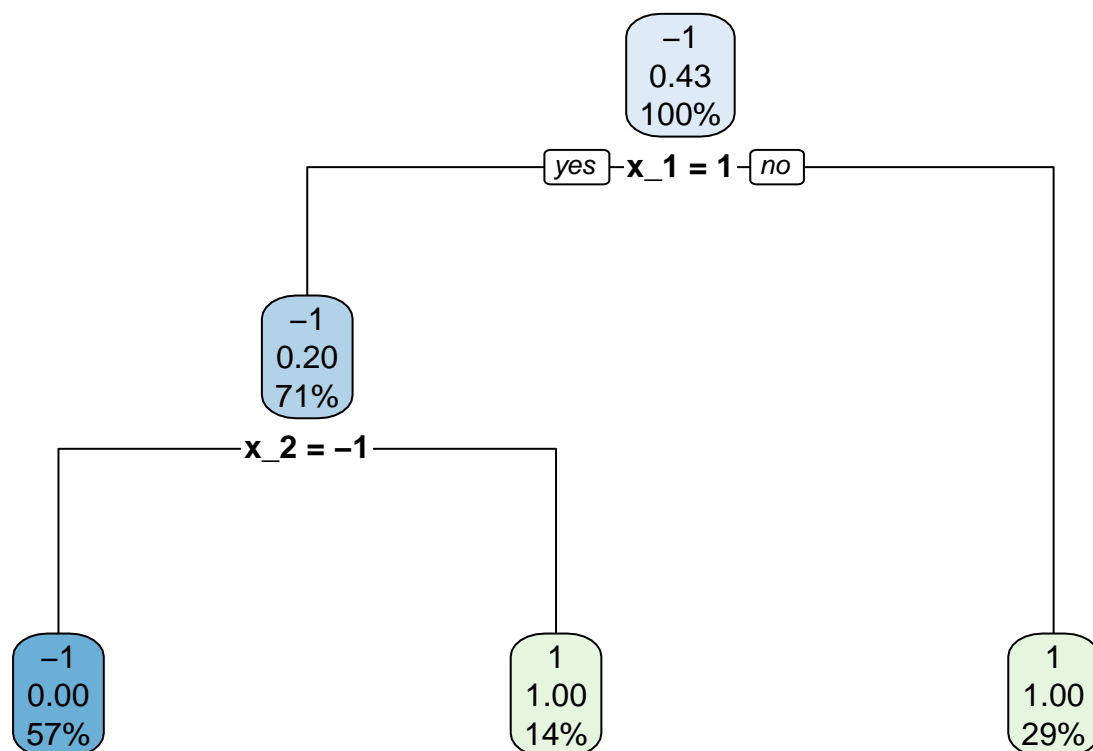
```
#split the decision tree based on gini index
shots = shots %>%
  mutate(x_1 = ifelse(x1>0.1, 1, -1) %>% as.factor(),
         x_2 = ifelse(x2>0.85, 1, -1) %>% as.factor())
fit1 = rpart(as.factor(y)~x_1+x_2, data = shots, method = "class", parms = list(split = "gini"),
            control=rpart.control(minsplit=1))
summary(fit1)
```

```
## Call:
## rpart(formula = as.factor(y) ~ x_1 + x_2, data = shots, method = "class",
##       parms = list(split = "gini"), control = rpart.control(minsplit = 1))
## n= 7
##
##          CP nsplit rel error   xerror   xstd
## 1 0.6666667     0 1.0000000 1.0000000 0.4364358
## 2 0.3333333     1 0.3333333 1.0000000 0.4364358
## 3 0.0100000     2 0.0000000 0.3333333 0.3086067
##
## Variable importance
## x_1 x_2
```

```

## 53 47
##
## Node number 1: 7 observations,    complexity param=0.6666667
## predicted class=-1 expected loss=0.4285714 P(node) =1
## class counts:      4      3
## probabilities: 0.571 0.429
## left son=2 (5 obs) right son=3 (2 obs)
## Primary splits:
## x_1 splits as RL, improve=1.8285710, (0 missing)
## x_2 splits as LR, improve=0.7619048, (0 missing)
##
## Node number 2: 5 observations,    complexity param=0.3333333
## predicted class=-1 expected loss=0.2 P(node) =0.7142857
## class counts:      4      1
## probabilities: 0.800 0.200
## left son=4 (4 obs) right son=5 (1 obs)
## Primary splits:
## x_2 splits as LR, improve=1.6, (0 missing)
##
## Node number 3: 2 observations
## predicted class=1 expected loss=0 P(node) =0.2857143
## class counts:      0      2
## probabilities: 0.000 1.000
##
## Node number 4: 4 observations
## predicted class=-1 expected loss=0 P(node) =0.5714286
## class counts:      4      0
## probabilities: 1.000 0.000
##
## Node number 5: 1 observations
## predicted class=1 expected loss=0 P(node) =0.1428571
## class counts:      0      1
## probabilities: 0.000 1.000
rpart.plot(fit1)


```



We can use “rpart” package to split our decision tree based on gini index. And this decision tree has all data points correctly classified.

```

ggplot(shots) + geom_point(aes(x = x1, y = x2, col = as.factor(y))) +
  theme_bw() + geom_vline(xintercept = 0.1, color = "red") + geom_hline(yintercept = 0.85, color = "blue") +
  stat_function(fun=function(x) sqrt(x), xlim = c(0,1))
  
```



We can choose  $x_1 > 0.1$  and  $x_2 > 0.85$  as our threshold for  $x_1, x_2$ , then split the decision tree based on transformed data. Using the reduction in the Gini index as the splitting criterion, the error is 0. We can adjust the threshold for  $x_1$  between  $[0.05, 0.15]$  and for  $x_2$  between  $[0.80, 0.95]$ , which will give us the same split on the data. The decision tree of such a threshold is plotted as follow.

```
shots = shots %>%
  mutate(x_1 = ifelse(x1>0.08, 1, -1) %>% as.factor(),
         x_2 = ifelse(x2>0.90, 1, -1) %>% as.factor())
fit2 = rpart(as.factor(y)~x_1+x_2, data = shots, method = "class", parms = list(split = "gini"),
            control=rpart.control(minsplit=1))
summary(fit2)
```

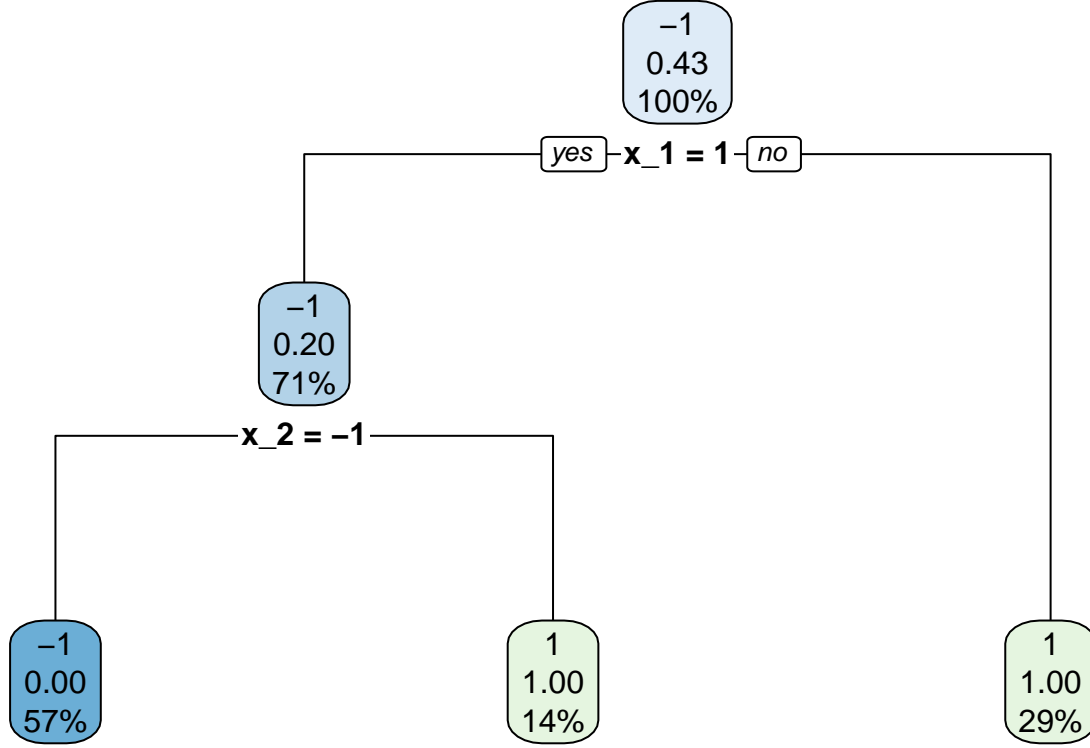
```
## Call:
## rpart(formula = as.factor(y) ~ x_1 + x_2, data = shots, method = "class",
##       parms = list(split = "gini"), control = rpart.control(minsplit = 1))
##      n= 7
##
##          CP nsplit rel error    xerror    xstd
## 1 0.6666667      0 1.0000000 1.0000000 0.4364358
## 2 0.3333333      1 0.3333333 1.0000000 0.4364358
## 3 0.0100000      2 0.0000000 0.3333333 0.3086067
##
## Variable importance
## x_1 x_2
## 53 47
##
## Node number 1: 7 observations,    complexity param=0.6666667
##   predicted class=-1 expected loss=0.4285714 P(node) =1
```

```

##      class counts:      4      3
##      probabilities: 0.571 0.429
##      left son=2 (5 obs) right son=3 (2 obs)
##      Primary splits:
##          x_1 splits as  RL, improve=1.8285710, (0 missing)
##          x_2 splits as  LR, improve=0.7619048, (0 missing)
##
## Node number 2: 5 observations,      complexity param=0.3333333
##      predicted class=-1 expected loss=0.2 P(node) =0.7142857
##      class counts:      4      1
##      probabilities: 0.800 0.200
##      left son=4 (4 obs) right son=5 (1 obs)
##      Primary splits:
##          x_2 splits as  LR, improve=1.6, (0 missing)
##
## Node number 3: 2 observations
##      predicted class=1 expected loss=0 P(node) =0.2857143
##      class counts:      0      2
##      probabilities: 0.000 1.000
##
## Node number 4: 4 observations
##      predicted class=-1 expected loss=0 P(node) =0.5714286
##      class counts:      4      0
##      probabilities: 1.000 0.000
##
## Node number 5: 1 observations
##      predicted class=1 expected loss=0 P(node) =0.1428571
##      class counts:      0      1
##      probabilities: 0.000 1.000

```

```
rpart.plot(fit2)
```



c

Since the three-point line is our approximation is  $x_2 = \sqrt{x_1}$ , we know  $f(x)^{\text{true}} = \text{sign}(x_2 - \sqrt{x_1})$ . We need to use a linear classifier that goes through origin, say  $f(\mathbf{x}) = \text{sign}(x_2 - ax_1)$ .

The true risk is  $R^{\text{true}}(f) = \mathbb{E}_{(x,y) \sim D} l(f(\mathbf{x}), y)$ , with the misclassification function  $l(f(\mathbf{x}), y) = 1_{\text{sign}(f(\mathbf{x}) \neq y)}$ .

Since  $x_1, x_2 \sim \text{Uniform}[0, 1]$ , we transform this problem into finding the definite integral of area between  $x_2 = \sqrt{x_1}$  and  $x_2 = ax_1$ , and it's intuitive that  $a \geq 1$ .

$$\begin{aligned}
 R^{\text{true}}(f) &= \int_0^{\frac{1}{a^2}} (\sqrt{x_1} - ax_1) dx_1 + \int_{\frac{1}{a^2}}^{\frac{1}{a}} (ax_1 - \sqrt{x_1}) dx_1 + \int_{\frac{1}{a}}^1 (1 - \sqrt{x_1}) dx_1 \\
 &= \frac{2}{3}a^{-3} - \frac{a}{2}a^{-4} + \frac{a}{2}a^{-2} - \frac{a}{2}a^{-4} - \frac{2}{3}a^{-\frac{3}{2}} + \frac{2}{3}a^{-3} + 1 - \frac{1}{a} - \frac{2}{3} + \frac{2}{3}a^{-\frac{3}{2}} \\
 &= \frac{1}{3}a^{-3} - \frac{1}{2}a^{-1} + \frac{1}{3}
 \end{aligned}$$

To minimize the true risk, we need to choose a value of  $a$  so that  $R^{\text{true}}(f)$  is minimized.


We have  $\frac{dR^{\text{true}}(f)}{da} = \frac{1}{2a^2} - \frac{1}{a^4}$ , which equals to 0 when  $a = \sqrt{2}$ . The second derivative shows that  $R^{\text{true}}(f)$  is increasing after  $a = \sqrt{2}$ . Hence,  $R^{\text{true}}(f)_{\min} = \frac{1}{3} - \frac{\sqrt{2}}{6}$ .

```

#plot observed data and another seperation line
ggplot(shots) + geom_point(aes(x = x1, y = x2, col = as.factor(y))) +
  theme_bw() + geom_abline(slope = -weight[100,1]/weight[100,2], color = 'red') +
  geom_abline(slope = sqrt(2), intercept = 0, color = 'blue') +
  stat_function(fun=function(x) sqrt(x), xlim = c(0,1))

```





The blue line shows the optimal linear classifier, but empirically it classified 1 point wrong. The empirical error  $R(f) = \frac{1}{8}$ . Any linear classifier with slope between  $(\frac{2}{3}, 5)$  would result in the same empirical error.

d

```
fit3 = rpart(as.factor(y)~x1+x2, data = shots, method = "class", parms = list(split = "gini"),
             control=rpart.control(minsplit=1))
summary(fit3)
```

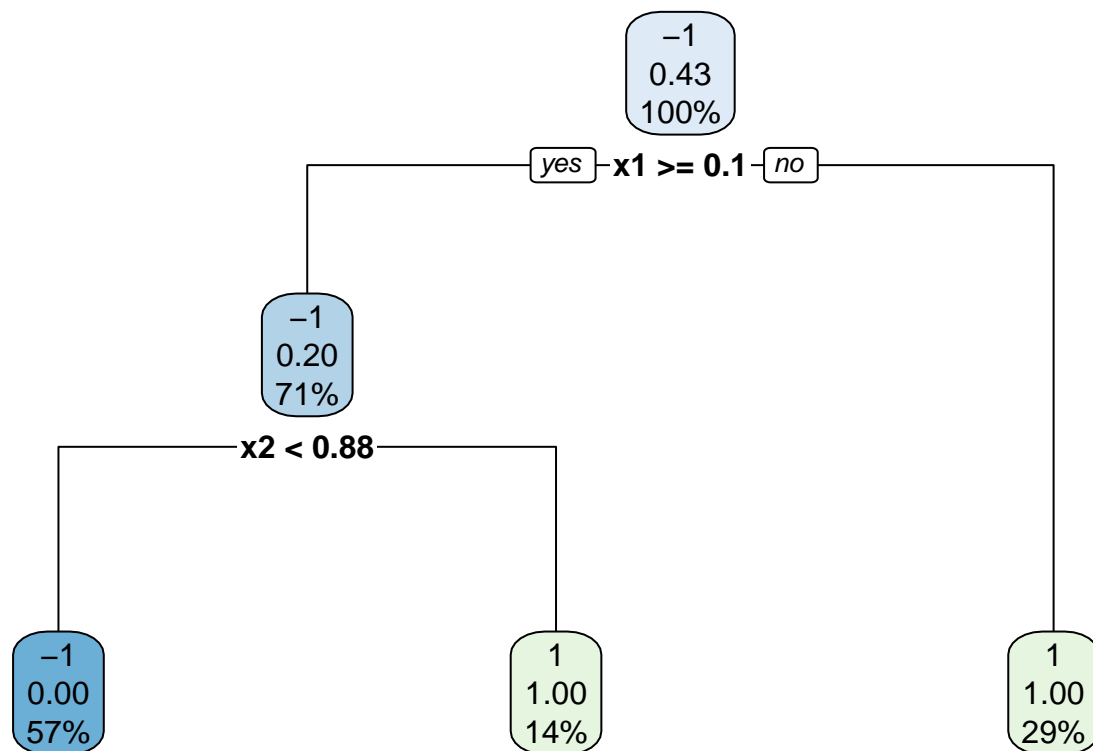
```
## Call:
## rpart(formula = as.factor(y) ~ x1 + x2, data = shots, method = "class",
##       parms = list(split = "gini"), control = rpart.control(minsplit = 1))
## n= 7
##
##          CP nsplit rel error  xerror    xstd
## 1 0.6666667      0 1.0000000 1.000000 0.4364358
## 2 0.3333333      1 0.3333333 1.333333 0.4364358
## 3 0.0100000      2 0.0000000 1.333333 0.4364358
##
## Variable importance
## x1 x2
## 53 47
##
## Node number 1: 7 observations,    complexity param=0.6666667
##   predicted class=-1 expected loss=0.4285714 P(node) =1
##   class counts:      4      3
```

```

##      probabilities: 0.571 0.429
##      left son=2 (5 obs) right son=3 (2 obs)
##      Primary splits:
##          x1 < 0.1   to the right, improve=1.828571, (0 missing)
##          x2 < 0.175 to the left,  improve=1.028571, (0 missing)
##
## Node number 2: 5 observations,      complexity param=0.3333333
##      predicted class=-1 expected loss=0.2 P(node) =0.7142857
##      class counts:      4      1
##      probabilities: 0.800 0.200
##      left son=4 (4 obs) right son=5 (1 obs)
##      Primary splits:
##          x2 < 0.875 to the left,  improve=1.6000000, (0 missing)
##          x1 < 0.8   to the left,  improve=0.2666667, (0 missing)
##
## Node number 3: 2 observations
##      predicted class=1 expected loss=0 P(node) =0.2857143
##      class counts:      0      2
##      probabilities: 0.000 1.000
##
## Node number 4: 4 observations
##      predicted class=-1 expected loss=0 P(node) =0.5714286
##      class counts:      4      0
##      probabilities: 1.000 0.000
##
## Node number 5: 1 observations
##      predicted class=1 expected loss=0 P(node) =0.1428571
##      class counts:      0      1
##      probabilities: 0.000 1.000

```


```
rpart.plot(fit3)
```



```

ggplot(shots) + geom_point(aes(x = x1, y = x2, col = as.factor(y))) +
  theme_bw() + geom_vline(xintercept = 0.1, color = "red") + geom_hline(yintercept = 0.88, color = "blue") +
  stat_function(fun=function(x) sqrt(x), xlim = c(0,1))

```



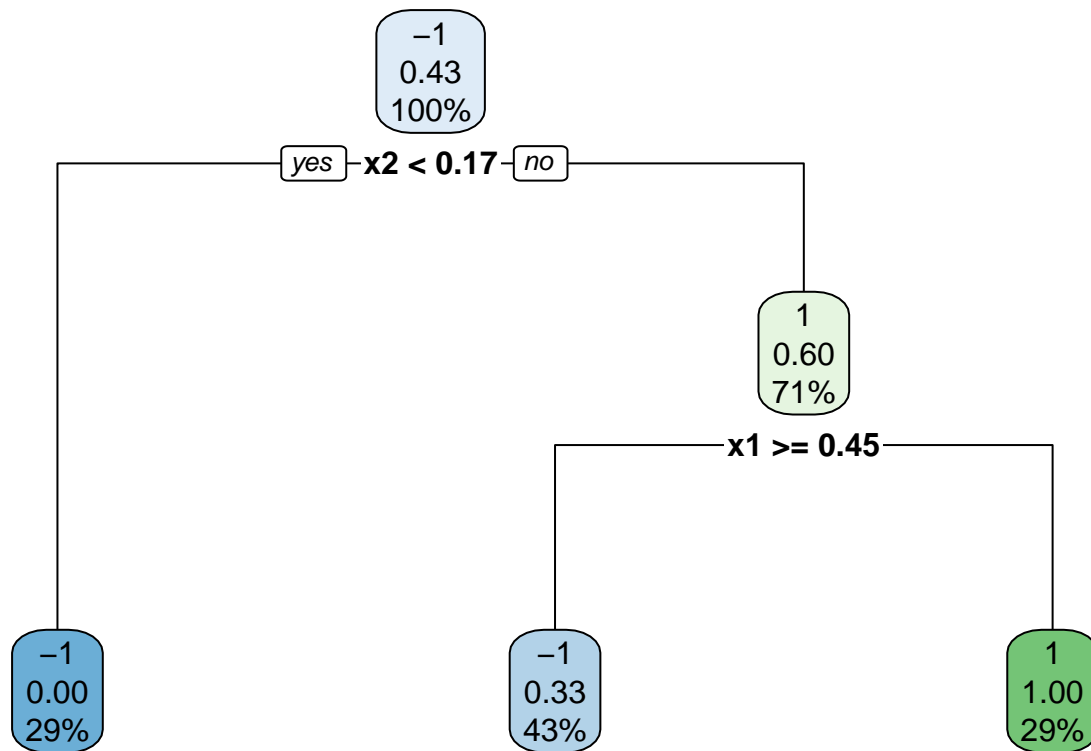
First decision tree splits on  $x_1$  first and the  $x_2$ , we can see  $s_1 = 0.1$  and  $s_3 = 0.88$ .  $s_2$  does not exist since there is no split with  $x_1 \leq 0.1$ . This is among the solutions that achieved the minimum empirical error.

```
#split on x2 first by assigning higher cost to x1
fit4 = rpart(as.factor(y)~x1+x2, data = shots, method = "class", parms = list(split = "gini"),
             control=rpart.control(minsplit=1, maxdepth = 2), cost = c(2,1))
summary(fit4)
```

```
## Call:
## rpart(formula = as.factor(y) ~ x1 + x2, data = shots, method = "class",
##       parms = list(split = "gini"), control = rpart.control(minsplit = 1,
##         maxdepth = 2), cost = c(2, 1))
##   n= 7
##
##           CP nsplit rel error   xerror   xstd
## 1 0.3333333      0 1.0000000 1.000000 0.4364358
## 2 0.0100000      2 0.3333333 1.666667 0.3984095
##
## Variable importance
## x2 x1
## 71 29
##
## Node number 1: 7 observations,      complexity param=0.3333333
##   predicted class=-1  expected loss=0.4285714  P(node) =1
##   class counts:      4      3
##   probabilities: 0.571 0.429
##   left son=2 (2 obs) right son=3 (5 obs)
##   Primary splits:
##     x2 < 0.175 to the left,  improve=1.0285710, (0 missing)
##     x1 < 0.1   to the right, improve=0.9142857, (0 missing)
##
## Node number 2: 2 observations
##   predicted class=-1  expected loss=0  P(node) =0.2857143
##   class counts:      2      0
##   probabilities: 1.000 0.000
##
## Node number 3: 5 observations,      complexity param=0.3333333
##   predicted class=1   expected loss=0.4  P(node) =0.7142857
##   class counts:      2      3
##   probabilities: 0.400 0.600
##   left son=6 (3 obs) right son=7 (2 obs)
##   Primary splits:
##     x1 < 0.45  to the right, improve=0.5333333, (0 missing)
##     x2 < 0.875 to the left,  improve=0.4000000, (0 missing)
##   Surrogate splits:
##     x2 < 0.65  to the right, agree=0.8, adj=0.5, (0 split)
##
## Node number 6: 3 observations
##   predicted class=-1  expected loss=0.3333333  P(node) =0.4285714
##   class counts:      2      1
##   probabilities: 0.667 0.333
##
## Node number 7: 2 observations
##   predicted class=1   expected loss=0  P(node) =0.2857143
##   class counts:      0      2
```


```
## probabilities: 0.000 1.000
```

```
rpart.plot(fit4)
```



```
#plot second tree
```

```
ggplot(shots) + geom_point(aes(x = x1, y = x2, col = as.factor(y))) +  
  theme_bw() + geom_vline(xintercept = 0.45, color = "red") + geom_hline(yintercept = 0.17, color = "blue") +  
  stat_function(fun=function(x) sqrt(x), xlim = c(0,1))
```




Second decision tree splits on  $x_2$  first and then  $x_1$ , we can see  $s_1 = 0.17$  and  $s_2 = 0.45$ .  $s_3$  does not exist since there is no split with  $x_2 \leq 0.17$ . This is not among the solutions that achieved the minimum empirical error.

e

Since we know the true function  $f(\mathbf{x}) = \mathbb{I}_{x_2 - \sqrt{x_1} > 0}$ , we can make a transformation  $x_2 = \sqrt{x_1}$ . Then our optimal linear classifier will simply be  $y = \mathbb{I}_{x_2 - x_1 > 0}$  and its error is 0.

```
#plot with transformation on x1
shots = shots %>%
  mutate(x1t = sqrt(x1))
ggplot(shots) + geom_point(aes(x = x1t, y = x2, col = as.factor(y))) +
  theme_bw() + geom_abline(slope = 1, color = 'blue')
```



f

```
fit5 = rpart(as.factor(y)~x1t+x2, data = shots, method = "class", parms = list(split = "gini"),
             control=rpart.control(minsplit=1))
summary(fit5)
```

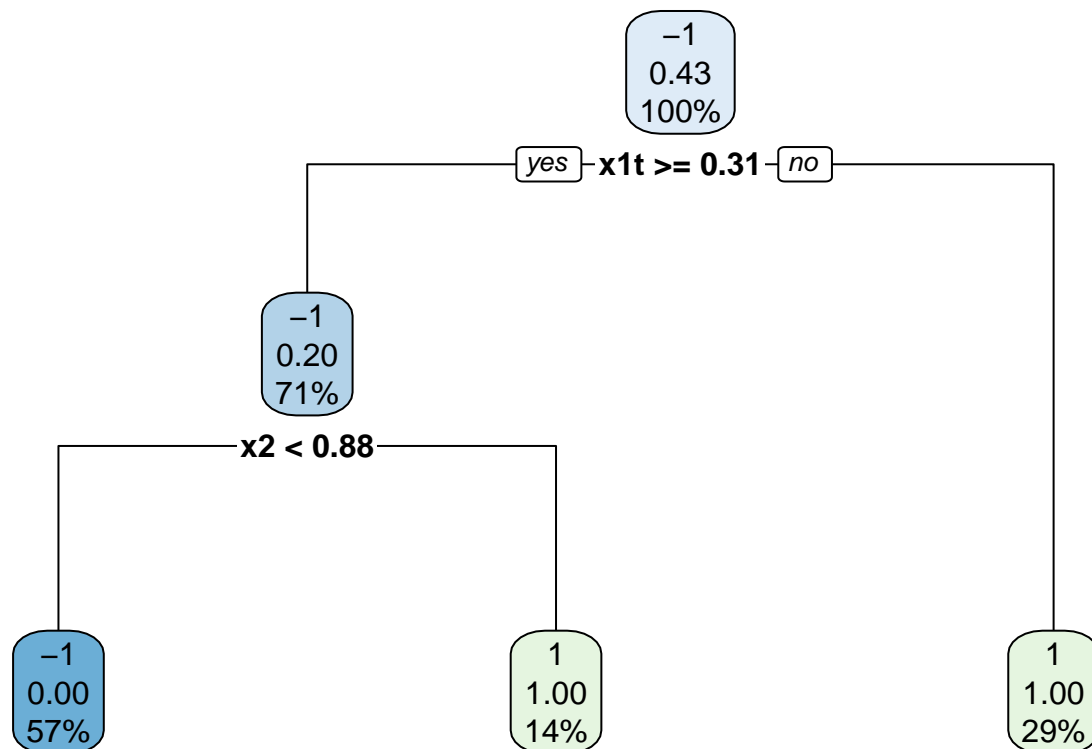
```
## Call:
## rpart(formula = as.factor(y) ~ x1t + x2, data = shots, method = "class",
##       parms = list(split = "gini"), control = rpart.control(minsplit = 1))
##      n= 7
##
##          CP nsplit rel error   xerror   xstd
## 1 0.6666667     0 1.000000 1.000000 0.4364358
## 2 0.3333333     1 0.333333 1.333333 0.4364358
## 3 0.0100000     2 0.000000 1.333333 0.4364358
##
## Variable importance
## x1t x2
## 53 47
##
## Node number 1: 7 observations,    complexity param=0.6666667
##   predicted class=-1 expected loss=0.4285714 P(node) =1
##   class counts:      4      3
##   probabilities: 0.571 0.429
##   left son=2 (5 obs) right son=3 (2 obs)
```

```

## Primary splits:
##   x1t < 0.3054526 to the right, improve=1.828571, (0 missing)
##   x2 < 0.175      to the left,  improve=1.028571, (0 missing)
##
## Node number 2: 5 observations,      complexity param=0.3333333
## predicted class=-1 expected loss=0.2 P(node) =0.7142857
##   class counts:      4      1
##   probabilities: 0.800 0.200
## left son=4 (4 obs) right son=5 (1 obs)
## Primary splits:
##   x2 < 0.875      to the left,  improve=1.6000000, (0 missing)
##   x1t < 0.8939899 to the left,  improve=0.2666667, (0 missing)
##
## Node number 3: 2 observations
## predicted class=1 expected loss=0 P(node) =0.2857143
##   class counts:      0      2
##   probabilities: 0.000 1.000
##
## Node number 4: 4 observations
## predicted class=-1 expected loss=0 P(node) =0.5714286
##   class counts:      4      0
##   probabilities: 1.000 0.000
##
## Node number 5: 1 observations
## predicted class=1 expected loss=0 P(node) =0.1428571
##   class counts:      0      1
##   probabilities: 0.000 1.000


```

```
rpart.plot(fit5)
```





```
ggplot(shots) + geom_point(aes(x = x1t, y = x2, col = as.factor(y))) +  
  theme_bw() + geom_vline(xintercept = 0.31, color = "red") + geom_hline(yintercept = 0.88, color = "blue")
```



After the transformation on  $x_1$ , the decision tree achieved the same error.

## h

We can represent the paint as area where  $x_1 \geq 0.5, x_2 \leq 0.25$ .

The true risk is  $R^{\text{true}}(f) = \mathbb{E}_{(x,y) \sim D} l(f(\mathbf{x}), y)$ , with the misclassification function  $l(f(\mathbf{x}), y) = 1_{\text{sign}(f(\mathbf{x}) \neq y)}$ .

Since  $x_1, x_2 \sim \text{Uniform}[0, 1]$ , we transform this problem into finding the definite integral of area between


$x_2 = \begin{cases} 0.25, & x_1 \geq 0.5 \\ 0, & x_1 < 0.5 \end{cases}$  and  $x_2 = ax_1$ , and it's intuitive that  $a \leq \frac{1}{2}$ .

$$\begin{aligned} R^{\text{true}}(f) &= \int_0^{\frac{1}{2}} ax_1 dx_1 + \int_{\frac{1}{2}}^{\frac{1}{4a}} \left(\frac{1}{4} - ax_1\right) dx_1 + \int_{\frac{1}{4a}}^1 (ax_1 - 1) dx_1 \\ &= \frac{3}{4}a + \frac{1}{4a} - \frac{9}{8} \end{aligned}$$

To minimize  $R^{\text{true}}(f)$ , we have  $\frac{dR^{\text{true}}(f)}{da} = \frac{3}{4} - \frac{1}{4a^2} = 0$ , which means  $a = \frac{\sqrt{3}}{3}$ . Since  $R^{\text{true}}(f)$  is decreasing on  $[\frac{1}{2}, \frac{\sqrt{3}}{3}]$  and increasing afterwards, we have  $R^{\text{true}}(f)_{\min} = \frac{\sqrt{3}}{2} - \frac{9}{8}$ .

i

The optimal decision tree will have decision boundaries that replicates the shape of paint and the error is 0.



2

a

i

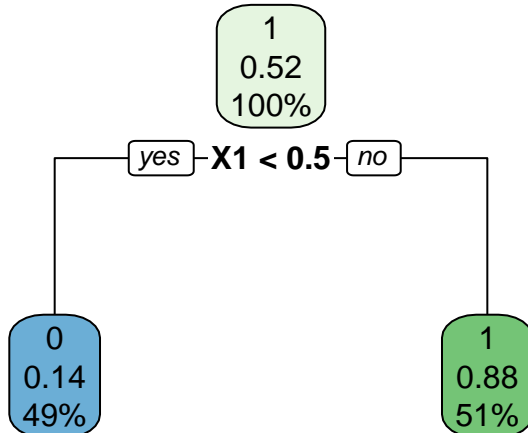
```
#read in datasets
train = read.csv("train.csv")
test = read.csv("test.csv")

#best split
stump1 = rpart(Y~., data = train, method = "class", parms = list(split = "gini"),
               control=rpart.control(minsplit=1, maxdepth = 1))
summary(stump1)

## Call:
## rpart(formula = Y ~ ., data = train, method = "class", parms = list(split = "gini"),
##       control = rpart.control(minsplit = 1, maxdepth = 1))
##       n= 500
##
##           CP nsplit rel error   xerror   xstd
## 1 0.7261411     0 1.0000000 1.0000000 0.04636138
## 2 0.0100000     1 0.2738589 0.2738589 0.03140616
##
## Variable importance
## X1 X2 X5
## 62 37  1
##
```

```
## Node number 1: 500 observations,      complexity param=0.7261411
##   predicted class=1  expected loss=0.482  P(node) =1
##   class counts:    241    259
##   probabilities: 0.482 0.518
##   left son=2 (243 obs) right son=3 (257 obs)
##   Primary splits:
##       X1 < 0.5 to the left,  improve=135.15930000, (0 missing)
##       X2 < 0.5 to the left,  improve= 52.78111000, (0 missing)
##       X3 < 0.5 to the left,  improve=  0.29823390, (0 missing)
##       X4 < 0.5 to the right, improve=  0.27751560, (0 missing)
##       X5 < 0.5 to the left,  improve=  0.05810746, (0 missing)
##   Surrogate splits:
##       X2 < 0.5 to the left,  agree=0.806, adj=0.601, (0 split)
##       X5 < 0.5 to the left,  agree=0.518, adj=0.008, (0 split)
##       X4 < 0.5 to the right, agree=0.516, adj=0.004, (0 split)
##
## Node number 2: 243 observations
##   predicted class=0  expected loss=0.1399177  P(node) =0.486
##   class counts:    209    34
##   probabilities: 0.860 0.140
##
## Node number 3: 257 observations
##   predicted class=1  expected loss=0.1245136  P(node) =0.514
##   class counts:     32   225
##   probabilities: 0.125 0.875
```

```
rpart.plot(stump1)
```



```
#best surrogate split
stump2 = rpart(Y~., data = train, method = "class", parms = list(split = "gini"),
               control=rpart.control(minsplit=1, maxdepth = 1, usesurrogate = 1))
summary(stump2)
```

```
## Call:
## rpart(formula = Y ~ ., data = train, method = "class", parms = list(split = "gini"),
##       control = rpart.control(minsplit = 1, maxdepth = 1, usesurrogate = 1))
##   n= 500
##
##           CP nsplit rel error    xerror    xstd
## 1 0.7261411     0 1.0000000 1.0000000 0.04636138
## 2 0.0100000     1 0.2738589 0.2738589 0.03140616
```

```
##
## Variable importance
## X1 X2 X5
## 62 37 1
##
## Node number 1: 500 observations,    complexity param=0.7261411
##   predicted class=1  expected loss=0.482  P(node) =1
##   class counts:    241    259
##   probabilities: 0.482 0.518
##   left son=2 (243 obs) right son=3 (257 obs)
##   Primary splits:
##     X1 < 0.5 to the left,  improve=135.15930000, (0 missing)
##     X2 < 0.5 to the left,  improve= 52.78111000, (0 missing)
##     X3 < 0.5 to the left,  improve=  0.29823390, (0 missing)
##     X4 < 0.5 to the right, improve=  0.27751560, (0 missing)
##     X5 < 0.5 to the left,  improve=  0.05810746, (0 missing)
##   Surrogate splits:
##     X2 < 0.5 to the left,  agree=0.806, adj=0.601, (0 split)
##     X5 < 0.5 to the left,  agree=0.518, adj=0.008, (0 split)
##     X4 < 0.5 to the right, agree=0.516, adj=0.004, (0 split)
##
## Node number 2: 243 observations
##   predicted class=0  expected loss=0.1399177  P(node) =0.486
##   class counts:    209    34
##   probabilities: 0.860 0.140
##
## Node number 3: 257 observations
##   predicted class=1  expected loss=0.1245136  P(node) =0.514
##   class counts:     32    225
##   probabilities: 0.125 0.875
```

```
rpart.plot(stump2)
```

